

Getting Started

[Updated: May 20, 2026]

Notes

To use SAKURA AI Engine, you will need a SAKURA internet member ID and a Sakura Cloud Project. For details, please refer to Sakura Cloud Startup Procedures.

- SAKURA AI Engine requires telephone authentication to confirm the customer's identity before starting use.
- To use SAKURA AI Engine, you can only pay by credit card, so you must register your credit card.

1. Account Signup

[Access the Control Panel.](#)

Agree to the SAKURA AI Engine Terms of Service.

Agreement to Terms of Service

Member ID

zyg31669

project

apitest

To use Sakura's AI Engine service, you must agree to [the Sakura AI Engine Terms of Service](#).

☐ I agree to the Terms of Service.

cancel

I agree to the Terms of Service.

Select the SAKURA AI Engine plan.

さくらのAI Engine

manual | | ? |

Plan

Available models

Account token

Document (RAG)

Usage amount

Playground

Plan > Change plan

Please select a plan.

Free base model plan

Basic fee 0 yen

✔ Chat generation (chat completions)

3000 requests per month

✔ Vector embeddings

10,000 requests per month

✔ Audio transcription

50 requests per month

✔ Audio speech

50 requests per month

• Rate limits will be applied to any amount exceeding the limit.

• Please note that even with the free base model plan, once you register documents for use with RAG, you will be continuously charged monthly based on the number of chunks until you delete the documents.

To sign a contract

Pay-as-you-go plan

Basic fee 0 yen

✔ Chat generation (chat completions)

3000 requests per month

✔ Vector embeddings

10,000 requests per month

✔ Audio transcription

50 requests per month

✔ Audio speech

50 requests per month

• Any excess usage will be charged on a per-use basis, depending on the number of input/output tokens used.

Cancel here

The basic model free plan can be used free of charge up to the following number of times, and if this is exceeded, a rate limit will be applied.

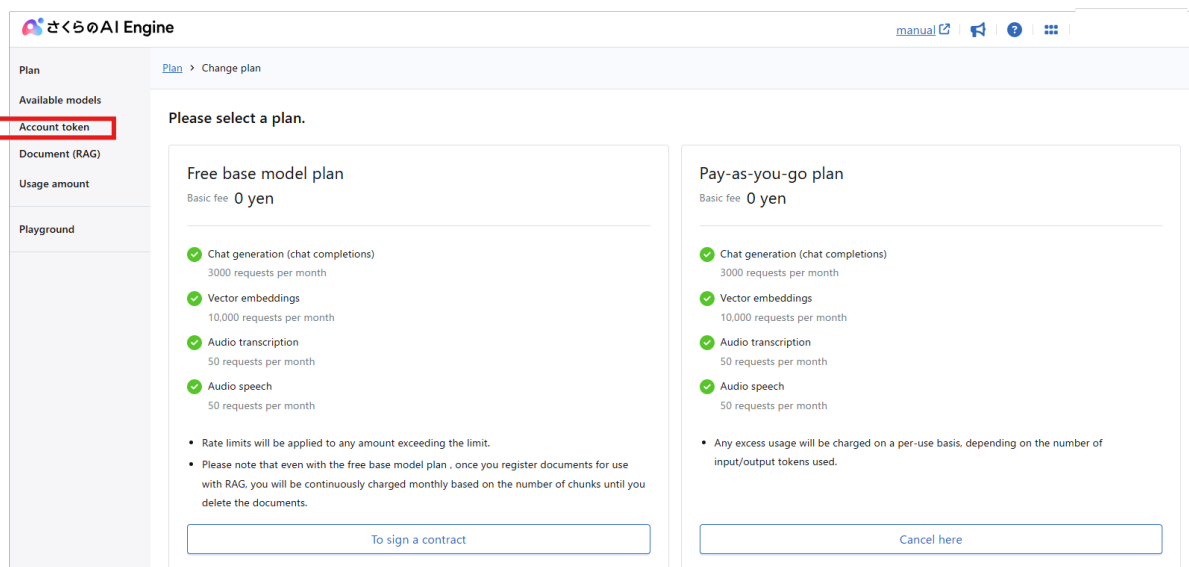
- Chat completions: 3,000 requests per month

- Vector embeddings: 10,000 requests per month
- Audio transcriptions: 50 requests per month
- TTS: 50 requests per month

2. Issue an account token

After selecting a plan, a menu will appear on the left side of the control panel.

Select `Account Tokens` from the left menu.



Click `Create Account Token` in the top right corner.



To get started using AI

To get started, first create an account token.

Create an account token

Enter the `Account Token Name` and click `Create`.

Create an account token

Please set up your account token.

Account token name

A name used to identify and manage account tokens.

cancel

Create

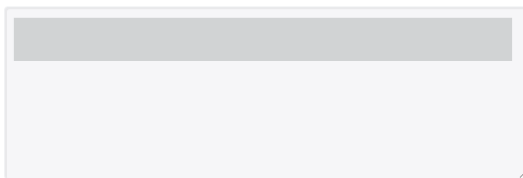
Your account token has now been issued.

Account token creation complete.

Your account token has been successfully created.

This token will not be displayed again. Please save it in a safe place.

Account token



copy

close

The account token is a ``<UUID>:<secret>`` pair. This token will not be displayed again. Copy it and save it in a safe place. If you lose your account token, issue a new token and delete the one you no longer need.

3. API endpoint

The API endpoints available with SAKURA AI Engine are as follows.

Chat & Messaging APIs

Used for text-based conversations and question-answering. Multimodal models also allow image input.

OpenAI compatible format:

- <https://api.ai.sakura.ad.jp/v1/chat/completions>
- <https://api.ai.sakura.ad.jp/v1/responses>

Anthropic compatible formats:

- <https://api.ai.sakura.ad.jp/v1/messages>

Vector Embeddings

Convert text to vector representation and use it for similarity searches, etc.

API endpoint: <https://api.ai.sakura.ad.jp/v1/embeddings>

Audio Transcription

Convert audio files into text.

API endpoint: <https://api.ai.sakura.ad.jp/v1/audio/transcriptions>

TTS

SAKURA AI Engine provides speech synthesis with the following endpoints.

OpenAI compatible format:

- <https://api.ai.sakura.ad.jp/v1/audio/speech>

VOICEVOX format:

- https://api.ai.sakura.ad.jp/tts/v1/audio_query
- <https://api.ai.sakura.ad.jp/tts/v1/synthesis>

4. Testing the token

Test using the issued account token.

Run the following curl command:

```
curl --location 'https://api.ai.sakura.ad.jp/v1/chat/completions' \
  --header 'Accept: application/json' \
  --header 'Authorization: Bearer <Token>' \
  --header 'Content-Type: application/json' \
  --data '{
    "model": "gpt-oss-120b",
    "messages": [
      { "role": "system", "content": "こんにちは！" }
    ],
```

```
"temperature": 0.7,  
"max_tokens": 200,  
"stream": false  
}'
```

Replace ``<Token>`` with your issued account token using the ``<UUID>:<secret>`` format, then run the command.

It is successful if you get a response like the one below.

Due to the technical characteristics of LLM, ``reasoning_content`` changes every time.

```
{  
  "id": "chatcpl-02b58c796a1c43aeb5893becc3a1d390",  
  "object": "chat.completion",  
  "created": 1757597188,  
  "model": "gpt-oss-120b",  
  "choices": [  
    {  
      "index": 0,  
      "message": {  
        "role": "assistant",  
        "content": "こんにちは！今日はどんなことをお手伝いしましょうか？😊",  
        "refusal": null,  
        "annotations": null,  
        "audio": null,  
        "function_call": null,  
        "tool_calls": [],  
        "reasoning_content": "We need to respond in Japanese presumably, as system says ¥"   
        こんにちは！¥". The user just said ¥"こんにちは！¥". We can respond friendly."  
      },  
      "logprobs": null,  
    }  
  ]  
}
```

```
    "finish_reason": "stop",
    "stop_reason": null
  }
],
"service_tier": null,
"system_fingerprint": null,
"usage": {
  "prompt_tokens": 73,
  "total_tokens": 127,
  "completion_tokens": 54,
  "prompt_tokens_details": null
},
"prompt_logprobs": null,
"kv_transfer_params": null
}
```

5. API documentation

API documentation for each endpoint can be found below.

[SAKURA AI Engine Inference API](#)