

User Guide

[Updated: July 15, 2026]

Available Models and API Documentation

Models provided by SAKURA AI Engine

The models that SAKURA AI Engine supports are as follows.

- Chat Completion Models
- Vector Embedding Models
- Audio Transcription Models

You can view the currently available models by selecting `Available Models` from the left menu in the Control Panel.

Available models (15)					Model licensing information
text audio					
For instructions on how to use the service, please refer to the user manual .					
Models with a status of "Application Required" are not available to individual customers. If you wish to use this service, please switch to a corporate or organizational account (not an individual account) and then apply.					
Chat generation (chat completions)					
Model Name	tag	status	Input fee	Output fee	
Qwen3-Coder-30B-A3B-Instruct	coding	End of offe			
Qwen3-Coder-480B-A35B-Instruct-FP8	coding	End of offe			
gpt-oss-120b		available	¥0.15 / 10,000 tokens (tax included)	¥0.75 / 10,000 tokens (tax included)	
llm-jp-3.1-b13b-instruct4	Made in Japan	available	¥0.15 / 10,000 tokens (tax included)	¥0.75 / 10,000 tokens (tax included)	
preview/Kimi-K2.6	Public Preview multimodal +3	available	¥0.60 / 10,000 tokens (tax included)	¥3.00 / 10,000 tokens (tax included)	
preview/Kimi-K2.7-Code	Public Preview multimodal +3	available	¥0.52 / 10,000 tokens (tax included)	¥5.04 / 10,000 tokens (tax included)	
preview/Phi-4-mini-instruct-cpu	Public Preview CPU	available	¥0.01 / 10,000 tokens (tax included)	¥0.03 / 10,000 tokens (tax included)	
preview/Phi-4-multimodal-instruct	Public Preview multimodal	End of offe			
preview/Qwen3-0.6B-cpu	Public Preview CPU	available	¥0.01 / 10,000 tokens (tax included)	¥0.03 / 10,000 tokens (tax included)	
preview/Qwen3-VL-30B-A3B-Instruct	Public Preview multimodal	available	¥0.10 / 10,000 tokens (tax included)	¥0.30 / 10,000 tokens (tax included)	
preview/Qwen3.6-35B-A3B	Public Preview multimodal +1	available	¥0.30 / 10,000 tokens (tax included)	¥1.50 / 10,000 tokens (tax included)	
preview/gemma-4-31B-it	Public Preview multimodal +3	available	¥0.24 / 10,000 tokens (tax included)	¥0.96 / 10,000 tokens (tax included)	
Vector embeddings					
Model Name	tag	status	Input fee		
multilingual-e5-large		available	¥2.00 / 10,000 tokens (tax included)		
preview/Qwen3-Embedding-4B-FP16	Public Preview	available	¥3.00 / 10,000 tokens (tax included)		
Audio transcription					
Model Name	tag	status	Input fee		
voice-large-v3-turbo		available	¥0.50 / 60 seconds (tax included)		

- TTS model

You can view the currently available models by selecting `Available Audio Models` from the left menu of the control panel.

text audio

For instructions on how to use the service, please refer to [the user manual](#). 

 Each voice model is provided as a "foundation model" in the "Sakura AI Engine Service Terms and Conditions."
Users must comply with the terms of use for each voice model, as stipulated in Article 9 of the "Sakura AI Engine Service Agreement."

Audio speech

Voice model name	style	Model ID	Voice ID	Speaker ID	status	Input fee 
Zundamon	normal ▼	zundamon	normal	3	 You need to check the term	—
Ankomon	normal ▼	ankomon	normal	113	 You need to check the term	—
Kasukabe Tsumugi	normal	kasukabetsumugi	normal	8	 You need to check the term	—
Himari Meimei	normal	meimeihimari	normal	14	 You need to check the term	—
Shikoku Metan	normal ▼	shikokumetan	normal	2	 You need to check the term	—
Tohoku Itako	normal	tohokuitako	normal	109	 You need to check the term	—
Tohoku Kiritan	normal	tohokukiritan	normal	108	 You need to check the term	—
Tohoku Zunko	normal	tohokuzunko	normal	107	 You need to check the term	—

API documentation

The API documentation is below.

- [Chat Completions / Audio Transcription / TTS](#)
- [Document / RAGAPI](#)

Chat Completions, Audio Transcription, Text-to-Speech, and the Document / RAG API use the token generated in the "Issue an Account Token" step.

Using the Document and RAG API [🔗](#)

SAKURA AI Engine provides an API and vector store for RAG in addition to the hosted LLM API.

Documentation/RAG API overview [🔗](#)

Retrieval-Augmented Generation (RAG) is a method in which a large-scale language model (LLM) searches (retrieval) an external knowledge base or document, and generates a response (generation) while incorporating the search results as input. This allows LLMs to provide information on internal company documents and the latest news that they would otherwise not know about.

Please refer to the Sakura Knowledge article for the RAG mechanism and processing flow.

When constructing a RAG, it is necessary to implement processes such as vectorizing external knowledge given in advance and storing it in a dedicated vector store, vectorizing questions submitted during a search, and performing a search in the vector store, but with SAKURA AI Engine, you can construct a RAG just by calling the API without implementing these processes.

With SAKURA AI Engine, you can vectorize pre-registered documents and store them in a vector store, search the vector store for questions, and convert them into natural language using the specified LLM with a single API call.

Document Upload API

You can add documents to the vector store using the document upload API (`documents_upload``). The uploaded document is vectorized with the specified embedding model and stored in the vector store.

APIs are also provided for document management, such as retrieving a list of documents, retrieving details, updating and deleting names and tags, and retrieving chunks. Please refer to [the RAG API specification](#) for details on each endpoint.

Supported Models for Document Upload

You can select from the following embedding models when uploading documents. The embedding model can be specified using the `model` parameter in your API request.

- `multilingual-e5-large` (default)
- `preview/Qwen3-Embedding-4B-FP16`

Chunk Size for Document Upload

Documents uploaded to SAKURA AI Engine are divided into specified token lengths, vectorized, and stored. This division unit is called a chunk, and its length is called the chunk size. Searches are performed chunk by chunk, and multiple chunks may be returned as search results. Chunk size can be specified with the ``chunk_size`` parameter when making API requests.

The chunk sizes that can be specified for each embedding model when uploading a document are as follows.

Model name	Minimum chunk size (token)	Maximum chunk size (token)	Default value (token)
<code>multilingual-e5-large</code>	256	512	512
<code>preview/Qwen3-Embedding-4B-FP16</code>	256	32000	512

Document Retrieval API

The following two types of APIs are available for referencing documents.

- `documents_query`: Performs a natural language search on the vector store. It outputs the

relevant documents or individual chunks stored in the vector store as-is.

- `documents_chat`: Performs natural language searches against the vector store and uses the search results to generate answers. Answers are generated in natural language by the specified LLM.

Normally, chat generation in RAG is done using `documents_chat` in three steps: 1. Vectorizing the submitted question 2. Searching the vector store 3. Converting the search results into natural language using the specified LLM, but SAKURA AI Engine processes this with a single API call.

When using the RAG API, usage charges apply for both the embedding model and the chat model.

Supported Distance Metrics for Vector Search

The following two types of distance metrics can be used in vector search. Distance metrics can be specified using the `distance_type` parameter when making API requests.

- `cosine`: Cosine distance
- `l2`: L2 distance (Euclidean distance)

RAG Quickstart

In this example, we will try RAG using a PDF file as our target document.

Instead of a simple full-text search across the ingested PDF, the LLM leverages vector search to generate natural, accurate responses.

Text import is compatible with common document formats (txt, pdf, html, docx, xlsx, md, etc.).

1. Vectorize the document and save it to the vector store

This is a process commonly called embedding. Prepare a suitable test PDF and run the following command. Replace `` with the account token obtained in the usage procedure. Replace it with the `` format and run it.

Document storage incurs fees. Please note that charges apply for document storage even under the Foundation Model Free Plan. Additionally, vectorizing documents incurs usage fees for the embedding model.

```
curl --request POST ¥
--url https://api.ai.sakura.ad.jp/v1/documents/upload/ ¥
--header 'Accept: application/json' ¥
--header 'Authorization: Bearer <Token>' ¥
--header 'Content-Type: multipart/form-data' ¥
--form "file=@document.pdf" ¥
--form "name=テストドキュメント" ¥
--form "tags=タグ 1" ¥
--form "tags=タグ 2" ¥
--form "model=multilingual-e5-large" ¥
```




```
--form "chunk_size=512"
```

Example response:

```
{ "id": "f9ccb16f-b231-45d6-a7db-7cdcd077638a", "status": "pending", "content": "", "name": "テストドキュメント", "tags": [ "タグ 1", "タグ 2" ], "model": "multilingual-e5-large", "chunk_size": 512 }
```

Wait for a while and select Document from the left menu of the console to see that the document has been imported.



When vectorizing documents, sentences are automatically extracted from uploaded files (PDF, text files, etc.) and used for searches and answers. However, depending on the file format and condition, the text may not be imported correctly.

Particularly in the following cases, accurate analysis may not be possible.

- PDF saved as an image (e.g. scan data)
- Special fonts and layouts (vertical writing, columns, heavy use of mathematical formulas and charts)

- Encrypted or password protected files
- Unsupported file format or corrupted file


If you are unable to import correctly, please try the following.

- Change the file format and upload again
- Check the status of the file you are uploading (is it corrupted, is it in the correct format, etc.)
- Convert it to a text file before uploading it

This service does not guarantee that the contents of all files can be completely imported.

2. Run a Search from the Control Panel

Enter the question you want to ask and `execute the query`, and the results will be output. We'll use `documents\_chat` to generate natural answers to our questions.



The screenshot shows the Sakura AI Engine Control Panel. The left sidebar contains links for Available models, Account token, Document (RAG) (selected), Usage amount, and Playground. The main area is titled 'Document (1)' and includes a search bar, a tag input field, and a 'search' button. A 'Query Search' button is also visible. Below the search bar is a table with columns: name, tag, Registration date, status, Number of documents, and Model.

name	tag	Registration date	status	Number of documents	Model
☐ test		2026/08/17 15:27:32	available	7	multilingual-e5-large

3. Run RAG via API

Run the following command:

```
curl --request POST ¥
--url https://api.ai.sakura.ad.jp/v1/documents/chat/ ¥
--header 'Accept: application/json' ¥
--header 'Authorization: Bearer <Token>' ¥
--header 'Content-Type: application/json' ¥
--data '{
  "model": "multilingual-e5-large",
  "chat_model": "gpt-oss-120b",
  "query": "質問文",
  "top_k": 3,
  "threshold": 0.3,
  "distance_type": "cosine",
}'
```

Using the Audio Transcription API

SAKURA AI Engine hosts a model called Whisper, which can transcribe audio data. It supports a wide range of audio and video formats, including mp3 and wav, as well as m4a and mp4, so you can use recorded data from meetings and interviews as is to create minutes and transcripts without having to convert the format in advance.

Audio transcription quick start

Below is an example of transcribing an MP3 audio file.

1. Request transcription of audio files

Because this API operates synchronously, audio files are limited to 30 minutes or 30 MB.

Run the following command: Replace `` with the account token obtained in the usage procedure. Replace it with the `` format and run it.

```
curl --request POST \
  --url https://api.ai.sakura.ad.jp/v1/audio/transcriptions \
  --header 'Accept: application/json' \
  --header 'Authorization: Bearer <Token>' \
  --header 'Content-Type: multipart/form-data' \
  --form 'file=@sample.mp3' \
  --form 'model=whisper-large-v3-turbo'
```

A response similar to the following will be returned.

```
{"text": "Hello. 皆さん聞こえますか", "model": "whisper-large-v3-turbo"}
```

Supported extensions

It supports common audio and video formats such as aac, m4a, mp3, mp4,

ogg, and wav (the list is just an example).

However, depending on the state of the file, it may not be possible to transcribe it correctly. If the process does not work correctly, please try the following.

- Change the audio file format and upload it again
- Check the status of the audio file you are uploading (is it corrupted, is it in the correct format, etc.)
- Convert to monaural audio and then upload

Additionally, this service does not guarantee that the contents of all audio files can be completely transcribed. The supported file formats listed are just examples, and the supported file formats are subject to change in the future.

How to use the TTS API

SAKURA AI Engine hosts speech synthesis models and can convert text to speech.

Notes

The speech synthesis API requires you to agree to the terms of use for each

speech synthesis model. Please check the terms of use for each speech synthesis model from the control panel.

TTS Quickstart

Below is a sample of converting text to speech.

1. Request Text-to-Speech Generation

Because this API operates synchronously, there is a limit of approximately 1,000 characters per request.

Run the following command: Replace ``<Token>`` with the account token obtained in the usage procedure. Replace it with the ``<uuid:secret>`` format and run it.

```
curl --request POST ¥
  --url https://api.ai.sakura.ad.jp/v1/audio/speech ¥
  --header 'accept: audio/wav' ¥
  --header 'Authorization: Bearer <Token>' ¥
  --header 'Content-Type: application/json' ¥
  --data '{
    "model": "zundamon",
    "input": "こんにちは、これは TTS のサンプルです。",
    "voice": "normal",
    "response_format": "wav"
  }' --output ./audio-speech-output.wav
```

In the sample request, the generated audio file is saved as `audio-speech-output.wav`.


Plan Cancellation and Handling of RAG Data Upon Cancellation ☒

No data is stored other than the documents you upload to RAG. When canceling, please delete all documents. Please note that you cannot cancel your subscription unless you delete all documents.

Cancellation Procedure ☒

1. Delete document ☒


Click `Delete` from the document details screen in the control panel.



Deletion can also be performed via the API. Please see the API documentation for details.

2. Delete account token

Delete the token by selecting Account Token from the left menu in the control panel.



3. Cancel Your Subscription

Click "Click here to cancel" on the plan selection screen, then click "Cancel plan" on the plan cancellation screen.

