

SAKURA AI Engine Inference API (1.0.0)

Download OpenAPI specification:[Download OpenAPI Spec\(YAML\)](#)[Download OpenAPI Spec\(JSON\)](#)

Copyright SAKURA internet Inc.

「 AI Engine Chat Generation / r e e c h a n ' c r i r t i o n / ' r e e c h a n t h e ' i ' e n d r o i n t

Basic usage

- [SAKURA AI Engine Getting Started | Sakura Cloud Manual](#)

API endpoint

<https://api.ai.sakura.ad.jp>

Create Chat Conversation

This is a typical example of a chat generation request. Please note that some parameters may not be supported depending on the model you use.

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: application/json

REQUIRED

model	string
required	Chat model name to use. Please check the control panel etc. for available models.
messages	Array of Developer message (object) or System message (object) or User message (object) or Assistant message (object) or Tool message (object)
required	Chat message history. The message types supported vary depending on the model.

max_tokens	integer >= 1 Maximum number of tokens to use for response generation.
temperature	number [0 .. 2] Default: 1 Parameter that controls generation diversity. Higher values produce more diverse responses, and lower values produce more deterministic responses.
tool_choice	string or ChatCompletionNamedToolChoice (object) (ChatCompletionToolChoiceOption)
tools	Array of ChatCompletionTool (object) List of tools available for the model.
stream	boolean Default: false Whether to enable streaming responses.

Responses

200

Success

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

POST /v1/chat/completions

Request samples

Payload

Content-Type: application/json

```
{
```

```
  "model": "string",
  "messages": [
    {}
  ],
  "max_tokens": 1,
  "temperature": 1,
  "tool_choice": "none",
  "tools": [
    {}
  ],
  "stream": false
}
```

}

Create a request

OpenAI compatible Responses API endpoint.

Please note that some parameters may not be supported depending on the model you use.

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: application/json

REQUIRED

model	string
required	Model name to use
input	string or Array of objects
required	Input to the model. Can be a string or an array of message objects

instructions	string System prompts inserted into the model context
stream	boolean Default: <code>false</code> Whether to enable streaming responses
temperature	number [0 .. 1] Parameter that controls generation diversity. Higher values produce more diverse responses, and lower values produce more deterministic responses.
top_p	number [0 .. 1] Nucleus sampling probability mass
max_output_tokens	integer Maximum number of output tokens to generate
tools	any List of tools that the model can call
tool_choice	any Control how the model uses tools
parallel_tool_calls	boolean Whether to allow parallel tool calls
store	boolean Whether to save the response of the generated model (item for OpenAI API compatibility). *Currently, it can be specified, but it will not be referenced on the system.
previous_response_id	string Previous response ID to chain (item for OpenAI API compatibility). *Currently, it can be specified, but it will not be referenced on the system.

conversation_id	string
	ID of the conversation to which this response belongs (item for OpenAI API compatibility). *Currently, it can be specified, but it will not be referenced on the system.
truncation	string
	Enum: "auto" "disabled"
	Model response truncation strategy
user	string
	Unique identifier representing the end user

Responses

200

Success

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

RO' T/v1/responses

Request samples

Payload

ÇöñtēñtÄyr ě

ă r lıçăŧiön/j' ôñ

{

"model": "string",

```
"input": "string",
"instructions": "string",
"stream": false,
"temperature": 1,
"top_p": 1,
"max_output_tokens": 0,
"tools": null,
"tool_choice": null,
"parallel_tool_calls": true,
"store": true,
"previous_response_id": "string",
"conversation_id": "string",
"truncation": "auto",
"user": "string"
```

}

Crăţăşă e' 'ăgă

Anthropic compatible Messages API endpoint.

Please note that some parameters may not be supported depending on the model you use.

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: application/json

REQUIRED

model	string
required	Model name to use
messages	Array of objects non-empty
required	List of message objects that make up the conversation
max_tokens	integer
required	Maximum number of tokens to generate
system	string

	system prompt
stream	boolean Default: false Whether to enable streaming responses
temperature	number [0 .. 1] Parameter that controls generation diversity. Higher values produce more diverse responses, and lower values produce more deterministic responses.
top_p	number [0 .. 1] Nucleus sampling probability mass
top_k	integer Top-k sampling parameters
stop_sequences	Array of strings List of sequences whose generation should be stopped
tools	Array of objects List of tools that the model can call
tool_choice	object or object Control how the model uses tools
metadata	object An optional key-value pair to attach to the request

Responses

200

Success

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

POST /v1/messages

Request samples

Payload

Content-Type: application/json

Request body

```
{
  "model": "string",
  "messages": [
    {}
  ],
  "max_tokens": 0,
  "system": "string",
  "stream": false,
  "temperature": 1,
  "top_p": 1,
  "top_k": 0,
  "stop_sequences": [
    "string"
  ],
  "tools": [
    {}
  ],
  "tool_choice": {
    "type": "auto",
    "disable_parallel_tool_use": true
  },
  "metadata": {
    "user_id": "string"
  }
}
```



```
}

```

Creare un API endpoint

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: multipart/form-data

REQUIRED

file	string <binary>
required	Audio file to transcribe. Common formats: aac, m4a, mp3, mp4, ogg, wav etc.
model	string Value: "whisper-large-v3-turbo" Transcription model identifier served by vLLM.
language	string Default: "ja" Source language hint (BCP-47, e.g. "ja", "en-US").
prompt	string Optional decoding/prompt bias (proper nouns, style hints).
temperature	number [0 .. 1] Default: 0 Decoding temperature.
stream	boolean Default: false

Responses

200
Transcription created

RESPONSE SCHEMA: application/json

model	string
required	Model used.
text	string
required	Full transcription (plain text).

- 400 Bad request
- 401 Unauthorized
- 429 Rate limited
- 500 Server error
- 504 Server error

POST /v1/audio/transcriptions

Response samples

200

```
{
  "model": "whisper-large-v3-turbo",
  "text": "Thank you for using our service today."
}
```

Create an embedding

This is a typical example of a request to create an embedding vector. Please note that some parameters may not be supported depending on the model you use.

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: application/json

REQUIRED

model	string
required	Name of the embedding model to use. Please check the Control Panel or relevant documentation for available models.
input	string or Array of strings
required	

Responses

200

Embedding created

RESPONSE SCHEMA: application/json

model	string
required	Model used.
data	Array of objects
required	

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

Request samples

Payload

Content-Type

application/json

```
{
  "model": "string",
  "input": "string"
}
```

Response samples

200

Content-Type

application/json

```
{
  "model": "multilingual-e5-large",
  "data": [
    {}
  ]
}
```

Content-Type

Generate speech from text (TTS).

- Required: input, model
- instructions can be specified but are currently ignored
- response_format can be specified, but currently it always returns wav
- stream is not supported (streaming is not possible even if stream_format is specified)

AUTHORIZATIONS:

BearerAuth

REQUEST BODY SCHEMA: application/json

REQUIRED

model

string

<code>required</code>	Speech synthesis model identifier (e.g. zundamon) Please check the control panel etc. for available models.
<code>input</code>	string [1 .. 1000] characters
<code>required</code>	Text to be synthesized (up to 1000 characters)
<code>voice</code>	string Speaker/Style (e.g. normal) Please check the control panel etc. for available voices.
<code>instructions</code>	string Additional instructions (e.g. tone of speech). *Currently, it can be specified, but it will be ignored.
<code>response_format</code>	string Default: "wav" Enum: "wav" "mp3" "ogg" "aac" "flac" *Currently, it can be specified, but wav is always returned.
<code>stream_format</code>	string Enum: "sse" "jsonl" *Currently, it can be specified, but it will be ignored.

Responses

200

Speech audio generated (always wav)

Response Headers

<code>Content-Type</code>	string Example: "audio/wav" Always <code>audio/wav</code> .
---------------------------	---

RESPONSE SCHEMA: audio/wav

string <binary>

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

RÔ' T/v1/audio/speech

Request samples

Payload

ÇõñtẽñțĂȳr ẽ

ă r lĩçăȳiõñ/j' ôñ

{

```
"model": "zundamon",
"voice": "normal",
"input": "Hello.",
"instructions": "speak in a calm tone",
"response_format": "wav",
"stream_format": "sse"
```

}

ÇrẽățẽĂũđĩõĂ ũěřȳĂT')

Create a query (JSON) for speech synthesis. Typically, you create a query in /tts/v1/audio_query and pass it to /tts/v1/synthesis to generate audio (wav). This API provides a compatible interface based on the /audio_query specification of the VOICEVOX Engine API.

Official specifications: https://voicevox.github.io/voicevox_engine/api/

AUTHORIZATIONS:

BearerAuth

QUERY PARAMETERS

text	string [1 .. 1000] characters
required	Text to synthesize into speech
speaker	integer >= 0
required	Speaker/Style ID (Please check the control panel etc. for available values)
enable_katakana_english	boolean
	Default: true
	Enable Katakana English.
core_version	string
	Version specification for speech synthesis.
	*Currently, it can be specified, but it will be ignored.

Responses

200

Audio query created

RESPONSE SCHEMA: application/json

accent_phrases	Array of objects (TtsAccentPhrase)
required	list of accent phrases
speedScale	number
required	overall speaking speed
pitchScale	number

<code>required</code>	overall pitch
<code>intonationScale</code>	number
<code>required</code>	overall intonation
<code>volumeScale</code>	number
<code>required</code>	overall volume
<code>prePhonemeLength</code>	number
<code>required</code>	Silence time before audio
<code>postPhonemeLength</code>	number
<code>required</code>	Silence after audio
<code>pauseLength</code>	number or null Silent periods such as punctuation marks. Ignored if null. Default value is null
<code>pauseLengthScale</code>	number Default: 1 Silence duration for punctuation marks (multiplier)
<code>outputSamplingRate</code>	integer
<code>required</code>	Audio data output sampling rate
<code>outputStereo</code>	boolean
<code>required</code>	Whether to output audio data in stereo
<code>kana</code>	string Reading (kana). [read-only] Text in AquesTalk style notation. Ignored as a query for speech synthesis
400 Bad request	
401	

Unauthorized

429

Rate limited

500

Server error

504

Server error

POST /tts/v1/audio_query

Response samples

200

Content-Type

application/json

{

```
"accent_phrases": [ ],
"speedScale": 1,
"pitchScale": 0,
"intonationScale": 1,
"volumeScale": 1,
"prePhonemeLength": 0.1,
"postPhonemeLength": 0.1,
"outputSamplingRate": 24000,
"outputStereo": false,
"kana": ""
```

}

Content-Type: application/json

Performs speech synthesis.

Generate audio (wav) by passing the query (JSON) created with /tts/v1/audio_query to the request body.

This API provides a compatible interface based on the /synthesis specification of the VOICEVOX Engine API.

Official specifications: https://voicevox.github.io/voicevox_engine/api/

AUTHORIZATIONS:

BearerAuth

QUERY PARAMETERS

speaker	integer ≥ 0
required	Speaker/Style ID (Please check the control panel etc. for available values)
enable_interrogative_upspeak	boolean
	Default: true
	Automatically adjust sentence endings when given interrogative text.
core_version	string
	Core Version. *Currently, it can be specified, but it will be ignored.

REQUEST BODY SCHEMA: application/json

REQUIRED

accent_phrases	Array of objects
required	list of accent phrases
speedScale	number
required	overall speaking speed
pitchScale	number
required	overall pitch
intonationScale	number
required	overall intonation
volumeScale	number
required	overall volume
prePhonemeLength	number

<code>required</code>	Silence time before audio
<code>postPhonemeLength</code>	number
<code>required</code>	Silence after audio
<code>pauseLength</code>	number or null Silent periods such as punctuation marks. Ignored if null. Default value is null
<code>pauseLengthScale</code>	number Default: 1 Silence duration for punctuation marks (multiplier). Default value is 1.
<code>outputSamplingRate</code>	integer
<code>required</code>	Audio data output sampling rate
<code>outputStereo</code>	boolean
<code>required</code>	Whether to output audio data in stereo
<code>kana</code>	string
<code>required</code>	Reading (kana). [read-only] Text in AquesTalk style notation. Ignored as a query for speech synthesis
<code>property name*</code>	
<code>additional property</code>	any

Responses

200

Speech audio generated (wav)

Response Headers

Content-Type

string

Example: "audio/wav"

audio/wav

RESPONSE SCHEMA: audio/wav

string <binary>

400

Bad request

401

Unauthorized

429

Rate limited

500

Server error

504

Server error

POST /tts/v1/synthesis

Request samples

Payload

Content-Type

application/json

{

"accent_phrases": [],

"speedScale": 1,

"pitchScale": 0,

"intonationScale": 1,

"volumeScale": 1,

"prePhonemeLength": 0.1,

"postPhonemeLength": 0.1,

"pauseLength": null,

"pauseLengthScale": 1,

"outputSamplingRate": 24000,

"outputStereo": false,

```
"kana": "string"
```

```
}
```