

SAKURA AI Engine

[Updated: June , 2026]

A secure and flexible API platform ideal for developing generative AI applications

SAKURA AI Engine is an API service that allows you to easily build generative AI applications. We carefully select open models and Japanese-specific models from a large number of models. Compatibility with OpenAI and Anthropic APIs allows for smooth migration from existing environments. Additionally, your data is never used for learning purposes, so you can use it with confidence even in operations that require a high level of information security. All data processing is completed on servers in Japan and operated in accordance with Japanese law. We provide the optimal AI platform for business with an emphasis on reliability and safety.

Try it on Playground first!



You can try some of the AI models provided by "SAKURA AI Engine" for free. You can check the operation intuitively using the GUI, so please use it for performance verification before full-scale implementation. Playground can be used with Sakura ID or GitHub/X (formerly Twitter) account authentication. (To use "SAKURA AI Engine," you need a Sakura member ID and a cloud project.)

Models you can try on Playground

- gpt-oss-120b
- llm-jp-3.1
- Qwen3-Coder
- Qwen3
- Phi-4
- VOICEVOX

Try Playground for free

OpenAI / Anthropic compatible

The latest LLM can be used with a compatible API, and domestic models can also be used.

Free plan available

Can be used like a pay-as-you-go plan until the usage limit is reached.

No leakage of usage data

Completed at DC in Japan Usage data is not sent to model provider

Secure closed configuration

VPN/LGWAN/Dedicated network compatible

Examples of LLMs available

Supports text generation, classification, embedding, and voice recognition (whisper)

Chat completions

Audio transcription

Embeddings

Document (RAG)

[View API documentation](#)

Pricing

Category	Provided model	Foundation Model Free Tier: Free usage up to request limits	Pay-as-you-go plan If the request for the free plan is exceeded	
Chat completions	gpt-oss-120b	Up to 3,000 requests per month	Input	
			0.15 / 10,000	
	Qwen3-Coder-480B- A35B-Instruct-FP8		Output	
			0.75 yen / 10,000	
			Input	
			0.3 yen/ 10,000	
			Output	

Audio transcription	Qwen3-Coder-30B-A3B-Instruct		2.5yen/ 10,000
		Input	
		0.15yen / 10,000	
		Output	
		0.75yen / 10,000	
		Input	
	llm-jp-3.1-8x13b-instruct4	0.15yen / 10,000	
		Output	
		0.75yen / 10,000	
		whisper-large-v3-turbo	Up to 50 requests per month
Embeddings	multilingual-e5-large	Up to 10,000 requests per month	Input
			2yen / 10,000
Text-to-Speech	VOICEVOX:ずんだもん	Up to 50 requests per month	Output
			Free
	VOICEVOX:東北ずん子		3yen / 10,000 mora
	VOICEVOX:東北きりたん		
	VOICEVOX:東北イタコ		
	VOICEVOX:四国めたん		
VOICEVOX:あんこもん			
VOICEVOX:Meimei Himari			
VOICEVOX:Tsumugi Kasukabe			
Document (RAG)		3yen / 100 chunks	

The displayed price includes tax.

The basic model free plan and the pay-as-you-go plan are separate plans (you will not automatically switch to the pay-as-you-go plan).

Even in the pay-as-you-go plan, there is a free usage limit that is equivalent to the base model free plan. There is no free usage tier for documents (RAG), and for both plans, fees are charged for each minimum unit (100 chunks) of use.

If you exceed the free usage limit on the basic model free plan, rate control will be applied to API requests. If you exceed the free usage limit on a pay-as-you-go plan, you will be charged for the minimum unit (10,000 tokens/60 seconds) of each base model.

There is a limit on the number of applications for the basic model free plan, and if the limit is reached, we will stop accepting new uses.

Please check [here](#) for license information and terms of use regarding the VOICEVOX provision model.

Mora is a unit that represents the length of the sound of one kana character.

Closed model pricing

A closed model is an AI model provided by the provider for commercial use. Review and permission are required based on the model provider's usage policy, so application is required for use in our environment. Applications can be made from the control panel of SAKURA AI Engine. Please check the [manual site](#) for details. [。](#)

Category	Provision model/provider	Base model free plan	pay-as-you-go plan
Chat Completions	Provided model: PLaMo 2.0-31B Provided by: PFN	Not applicable	For more information, please check the manual site

Provided model: cotomi v3* Provided by: NEC	or contact us individually. Contact form
---	--

cotomi: "cotomi" is the name of the generative AI developed by NEC. The name reflects NEC's vision to illuminate the future through language—helping ideas flourish ("minoru") into reality ("koto")—and its commitment to being a dedicated AI partner for customers. ("cotomi" is a registered trademark of NEC Corporation.)

public preview model

Provided model	Base model free plan	pay-as-you-go plan
Qwen3-0.6B-cpu (from October 9, 2025~)	Up to 3,000 requests per month	Input 0.01yen / 10,000
		Output 0.03yen / 10,000
Phi-4-mini-instruct-cpu (from October 9, 2025~)		Input 0.01yen/ 10,000
		Output 0.03yen/ 10,000
Qwen3-VL-30B-A3B-Instruct (from October 21, 2025~)		Input 0.1yen/ 10,000
		Output 0.3yen/ 10,000
Phi-4-multimodal-instruct (from October 21, 2025~)		Input 0.1yen/ 10,000
		Output 0.3yen/ 10,000
Kimi-K2.6 (from May 12, 2026~)		Input 0.6yen/ 10,000
		Output 3yen/ 10,000
Qwen3.6-35B-A3B (from June 4, 2026~)		Input 0.3yen/ 10,000

gemma-4-31B-it (from June 30, 2026~)		Output
		1.5 yen/ 10,000
		Input
		0.24 yen/ 10,000
Qwen3-Embedding-4B(FP16) (from December 16, 2025~)	Up to 10 ,000 requests per month	Output
		0.96 yen/ 10,000
		Input
		3 yen/ 10,000
		Output
		Free

The public preview model is provided to users for the purpose of verifying operation in an actual usage environment and gathering feedback, and may be discontinued without prior notice. Even during the provision period, we do not guarantee stability or response quality, and operation and specifications may change. After the public preview ends, prices may change if we continue to provide the service.

Usage flow

STEP 1.

Create a SAKURA internet member ID

If you do not have a member ID, first create one. A phone number is required for registration. (Telephone number verification will be performed) If you already have a SAKURA internet member ID, please proceed to step 2.

Create a SAKURA internet member ID (free)

STEP 2.

Create a “Sakura Cloud” project

After logging in to the control panel using your member ID, create a Sakura Cloud project. Credit card registration is required to use "Sakura Cloud". If you have registered your credit card with other SAKURA internet services, no additional registration is required. If you already have a "Sakura Cloud" project, please proceed to step 3.

Create/Log in to "Sakura Cloud" project

STEP 3.

Start using from the control panel

Please access the control panel and use SAKURA AI Engine. You will be redirected to the Sakura Cloud login screen. Please enter the account information obtained in STEP1 and STEP2 and log in.

Access the “SAKURA AI Engine” control panel

Contact form