



Building a Web App with AI Engine

SAKURA AI Engine Practical Exercise Module4

<https://www.sakura.ad.jp/>

DATE

2026.9

COMPANY

SAKURA internet Inc.

SERVICE

SAKURA AI Engine

VERSION

Ver.1.0

Today's Agenda

SAKURA AI Engine Hands-On consists of four sessions, 1 through 4.

Module 1	SAKURA AI Engine Practical Exercise 1 Log in to AI Engine, select a model, issue an account token, run inference, and build a RAG chatbot fed with a document
Module 2	SAKURA AI Engine Practical Exercise 2 A setup that uses AI to automate everything from meeting-audio transcription to summarization and minutes
Module 3	SAKURA AI Engine Practical Exercise 3 Combining image and text input to design AI output that supports business decisions
Module 4	SAKURA AI Engine Practical Exercise 4 Embedding text AI calls into a Web app, all the way to running it in the browser

★ You are here

Goal of This Class

Use the distributed code to understand how a Web app built with AI Engine works.

What We will Do

- 1 Set up the runtime environment
- 2 Read through the code's flow
- 3 Fill in a few blanks in the code
- 4 Run it, ask a question, and check the answer and its sources

The App We will Build

The screenshot shows a web application titled "RAG質問アプリ". It contains a text input field for a question, with a sample question in Japanese: "伊勢物語における在原業平について登録ドキュメントの内容を教えてください". Below the input field is a section for "RAGパラメータ" (RAG Parameters). It includes a "top_k" input field set to "3", a "threshold" slider set to "0.3", and explanatory text in Japanese about how these parameters affect the search results.

- Send a natural-language question against the registered RAG data.
- View the result in your browser

- Change top_k / threshold from the screen
- Display the answer and its sources (name, chunk, distance, content)

Turn the RAG you ran with curl into something you can use from a browser


Overview 1

The Environment We'll Build

WSL + Python Virtual Env (venv) + Flask

The System We'll Build in This Class

The user is on the left, AI Engine on the right. The question goes out, and the answer comes back.



Outbound

Question "Which Thai restaurant is the most popular in Thailand?" $\text{top_k} = 3$ $\text{threshold} = 0.5$


Inbound

Answer "According to the registered documents, ●● and XX are popular in Bangkok."

Sources = [document name / chunk number / distance / content] for the top top_k results

Environment Overview

Set up a Windows + Linux, and inside that a Python virtual env, in your PC.




Set up your own Windows + Linux + Python environment for this class.


Overall Structure (Data Flow)

Until now you called curl directly with AI Engine; now call it through the browser and Flask.

Outbound (Request)



Reply (Inbound Response)



The account token is held only by Flask; it stays hidden from the browser and JavaScript.

Why Flask Is Needed

Browser and AI Engine need a Web server in between. That's Flask.

Entry Point

The window that receives requests from the browser (routing).

Mediates Between Browser and AI Engine

“requests” calls AI Engine and formats the result before returning it.

Hides the Token

Holds the account token on the backend, never exposing it to the browser.

A one-off curl isn't enough —Flask is needed to turn this into a Web app the screen can call any number of times.

What Is venv?

A separate, independent Python environment. Also called a virtual environment.

Why Separate Environments?

- Avoids conflicts between dependencies
- Doesn't affect other projects
- Installs only what you need
- Easy to reproduce the environment

Picture It

Install everything system-wide

Versions clash easily

venv separates things into boxes

Independent and safe, per project

This time we'll create a dedicated box with only Flask and requests installed.

Directory Structure

Extract the distributed zip and lay it out like this.

04_rag-web-app/

```
|— app.py
|— requirements.txt
|— README.md
|— templates/
|   |— index.html
|— static/
|   |— app.js
|   |— style.css
```

templates/

Where Flask's returned HTML lives

static/

Where the JS/CSS the browser reads lives

Flask auto-loads templates and static from fixed locations, so it must be laid out this way.


Overview 2

curl vs. Turning It Into a Web App

From calling curl directly, to an app with a screen.

How We've Done It So Far

In sessions 1–3, you used curl from the terminal to call AI Engine directly.



- Run commands by hand
- Only you run it
- Result shows in the terminal
- One-off calls, one at a time

It works, but it's all self-contained inside your own terminal.

Web App: The Challenge

curl is for things like checking behavior works; it's not something anyone can just use.

A Command Every Time

Users must remember the terminal and the command.

The Token Is Exposed

The token rides along in the command. Share it, and it leaks.


No Screen

No input field, no result display — you can't hand it to someone. You'd have to read the code and parse the result — a hassle

From “I run it myself” to “someone else can use it”.

What It Means to Turn This Into a Web App

A screen instead of a command; backend processing instead of manual steps.



With curl

Type a command

Shows in the terminal

Token on hand



With the Web App

Click a button on screen


Formatted and shown on screen

Hidden inside the server

By operating the browser, the user can trigger a call to AI Engine.

Why a Server Needs to Sit in Between

Anyone can view the browser's code, so secrets belong on the server side.




That's why a server in between —Flask — is needed.


Before / After

The task is the same. how you call it changes.

Before: Invoking via the curl command



After: Web App



Action: command → screen | Display: terminal → screen | Token: on hand → server

Who you call (AI Engine) is the same. What changes is which endpoint and how you call the API.

Hands-on 1

Environment Setup

WSL assumes it's already installed. From here, you'll actually start working hands-on.

Setup Flow

We'll proceed in 4 steps.

- 1 Place the distributed files in your WSL home
- 2 Unzip distributed files
- 3 Create and enter a virtual environment (venv)
- 4 Install packages and set the token

Work through them in order, and let's launch the Web app!

What's Inside, and What You'll Build

These are the 6 files you'll be working with, and the folder shape you'll build with them.

app.py

The Flask server itself. Receives requests, validates, calls requests, formats results

index.html

The screen's skeleton (HTML)

app.js

Gets input, sends fetch requests, shows results (frontend logic)

style.css

The screen's look (styling)

requirements.txt

List of required packages (Flask, requests)

README.md

Usage instructions

Target Folder Structure

```
04_rag-web-app /
├── app.py
├── requirements.txt
├── README.md
├── templates/
│   └── index.html
└── static/
    ├── app.js
    └── style.css
```

These 6 files complete the round trip: browser → Flask → AI Engine. Downloading and extracting next will arrange them into this shape automatically.

Download and Extract the Distributed Files

Download and extract the exercise files — all from your WSL terminal, no browser or Explorer needed.

```
$ cd ~  
$ wget https://github.com/onigiribouya/sakura-ai-engine-module4-download-  
zip/raw/refs/heads/main/module4-src.zip  
$ python3 -m zipfile -e module4-src.zip 04_rag-web-app
```

- Both commands run in your WSL terminal (the black window).
- No output after these commands? That's normal — it means they worked.
- If you run wget more than once, it may save as module4-src.zip.1 or .2 instead of overwriting. If extraction looks wrong, run `rm -f module4-src.zip*` and try again.

Now `~/04_rag-web-app` has all the distributed files, already laid out correctly for Flask.

Check the Folder Structure

Let's confirm it matches the shape from earlier.

Check It

```
$ cd ~/04_rag-web-app  
$ ls -l
```

- You should see app.py, requirements.txt, README.md, templates/, and static/.
- Inside templates/: index.html. Inside static/: app.js, style.css.

Target Folder Structure

```
04_rag-web-app /  
├─ app.py  
├─ requirements.txt  
├─ README.md  
├─ templates/  
│   └─ index.html  
└─ static/  
    ├─ app.js  
    └─ style.css
```

Everything's in place — let's set up the Python environment next.

Create a Virtual Environment (venv)

04_rag-web-app : create a dedicated Python environment inside it.

```
$ python3 -m venv venv
```

- **Inside** 04_rag-web-app, a folder called venv (your dedicated environment) is created.
- On WSL (Ubuntu), use python3 instead of python.
- You only create it once; no need to redo it after that.
- How to check it worked: run ls — you should see a new **venv** folder.

Example: ls -l output after venv is created

```
-rw-r--r-- 1 k- k- 10142 Aug 8 05:35 app.py
-rw-r--r-- 1 k- k- 15 Aug 5 02:18 requirements.txt
drwxr-xr-x 2 k- k- 4096 Aug 8 05:13 static
drwxr-xr-x 2 k- k- 4096 Aug 8 05:13 templates
drwxr-xr-x 5 k- k- 4096 Sep 8 22:49 venv
```

At this point, you haven't entered the environment yet. The next step enters it.

Entering / Leaving venv

The sign to watch for is **(venv)** at the start of the prompt.
With (venv) showing, you need to run Flask.

Entering the virtual env (activate)

```
user@host:~/04_rag-web-app$ source venv/bin/activate  
  
(venv) user@host:~/04_rag-web-app$
```

At the front, **(venv)** once it shows, you're in the virtual environment.

Leaving (deactivate)

```
(venv) user@host:~/ 04_rag-web-app$ deactivate  
  
user@host:~/04_rag-web-app $
```

At the front, **(venv)** once it's gone, you've left the virtual environment.

From here on, do all work with **(venv) showing.**

Install the Packages

(venv) is showing before running this.

```
(venv) $ pip install -r requirements.txt
```

- requirements.txt lists Flask and requests, which get installed.

- (venv) If you run this without showing, it installs elsewhere. Be sure to check.

- You can check whether it's installed with pip list

- Look for Flask and requests in the list — the other packages are their dependencies, installed automatically.

Now the environment to run the app is ready!

```
(venv) ~:~/04_rag-web-app$ pip list
```

Package	Version
blinker	1.9.0
certifi	2026.7.22
charset-normalizer	3.5.1
click	8.5.0
Flask	3.1.3
idna	3.19
itsdangerous	2.2.0
Jinja2	3.1.6
MarkupSafe	3.0.3
pip	24.0
requests	2.34.2
urllib3	2.7.0
Werkzeug	3.1.8

Set the Token

Set the account token as an environment variable.

```
(venv) $ export AI_ENGINE_TOKEN="<your token>"
```

- Flask reads the token from this environment variable. The browser and JS never get to see the token's contents.
- This setting disappears when you close that terminal. If you reopen it, set it again.
- You can check whether it's set with `echo $AI_ENGINE_TOKEN`.

Never write the token directly in code — read it from an environment variable.

Setup Complete Checklist

You're ready once these 4 are all true.

- ✓ `04_rag-web-app` has `templates/` and `static/` laid out correctly
- ✓ shows `(venv)` at the start of the prompt
- ✓ `pip list` has Flask and requests
- ✓ `echo $AI_ENGINE_TOKEN` shows the token

Next, move on to filling in the code and running it!

Hands-on 2

Filling In the Code

Think about where in the code the values you typed with curl should go

The Big Picture of the Blanks

3 groups, 7 blanks in total. **Build** → **Send** → **Receive** — that's the order we'll follow.



- 1 **Build what to send** `app.py build_payload (__(1)__ __(2)__ __(3)__)`
- 2 **Send what you built** `app.py requests.post (__(4)__ __(5)__)`
- 3 **Retrieve the result that comes back** `app.js`

Build the data to send, send it to AI Engine, receive the data that comes back, and display it. Think it through in this order.

1 Build What to Send

Put the values you passed via `--data` into payload below.

`app.py` `build_payload` (fill in the blanks)

```
def build_payload(query, top_k, threshold):
    payload = {
        "model": EMBEDDING_MODEL,
        "chat_model": CHAT_MODEL,
        "query": __ (1) __,
        "top_k": __ (2) __,
        "threshold": __ (3) __,
        "distance_type": DISTANCE_TYPE,
    }
```

The `--data` you saw with curl

```
--data '{
  "query": "...",
  "top_k": 3,
  "threshold": 0.5
}'
```

What goes in?

- `model` • `chat_model` • `distance_type` are fixed constants. They aren't blanks.

The 3 values from curl's `--data`. What goes in each one?

1 Answer: Try Plugging In Values First

Here's what it looks like if you write the values directly, the way they were set on the screen.

app.py build_payload (with values filled in)

```
payload = {  
    "model": EMBEDDING_MODEL,  
    "chat_model": CHAT_MODEL,  
    "query": "Which Thai restaurant is best?",  
    "top_k": 3,  
    "threshold": 0.5,  
    "distance_type": DISTANCE_TYPE,  
}
```

Screen Input (example)

Question

Can I return this?

top_k

3

threshold

0.5

As is, the value is fixed: whatever you type on the screen won't be reflected.

1 Answer: It's Actually in Variables

The values entered on the screen are in the variables `query` / `top_k` / `threshold`.

`app.py build_payload(answer)`

```
payload = {  
    "model": EMBEDDING_MODEL,  
    "chat_model": CHAT_MODEL,  
    "query": query,  
    "top_k": top_k,  
    "threshold": threshold,  
    "distance_type": DISTANCE_TYPE,  
}
```

Screen Input → Variable

The screen's "Question" → `query`

The screen's "Number of References" → `top_k`

The screen's "Threshold" → `threshold`

- The variable holds whatever value was entered on the screen at that time. In other words, writing the variable name means the input goes straight into payload.

Assign the variable, not a fixed value. The value entered on the screen goes straight to AI Engine.

1 Answer (Final Version)

Just take the value received as an argument and put it straight into payload.

app.py build_payload(answer)

```
payload = {  
    "model": EMBEDDING_MODEL,  
    "chat_model": CHAT_MODEL,  
    "query": query,  
    "top_k": top_k,  
    "threshold": threshold,  
    "distance_type": DISTANCE_TYPE,  
}
```

__(1)__ query

The question text (validated string)

__(2)__ top_k

Maximum number of documents to retrieve

__(3)__ threshold

Vector-distance threshold (0–2)

The same 3 from curl's --data. Even the names stay the same.

2 Send What You Built

Where do curl's destination URL and --data go inside requests.post?

app.py requests.post (fill in the blanks)

```
response = requests.post(  
    __ (4) __ ,  
    headers=headers,  
    json=__ (5) __ ,  
    timeout=REQUEST_TIMEOUT,  
)
```

What it looked like with curl

```
curl <destination URL> ¥  
  --header "... " ¥  
  --data '{...}'
```

What goes in?

- URL → 1st argument, --data → json=. That's the mapping.

Map curl's URL and --data. In requests, that's the 1st argument and json=.

2 Answer

Send the payload you built in 1 , here.

app.py `requests.post(answer)`

```
response = requests.post(  
    AI_ENGINE_URL,  
    headers=headers,  
    json=payload,  
    timeout=REQUEST_TIMEOUT,  
)
```

__(4)__ `AI_ENGINE_URL`

Destination (the constant defined above)

__(5)__ `payload`

The body you built in 1

- curl's URL is the 1st argument, --data is json=.

Send the payload you built in 1, here.

3 Retrieve the Result That Comes Back

app.js reads the JSON that Flask returns, and pulls out the answer and its sources.

Located in static/

app.js Extracting the response (fill in the blanks)

```
// 7. Display the answer and its sources
showAnswer(data.answer);
showSources(data.sources);
```

- Path: static/app.js

Shape of the JSON that comes back

```
{
  "answer": "...",
  "sources": [ {...}, {...} ]
}
```

What goes in?

- Write the JSON key name directly after the . dot

data.answer and **data.sources** map directly to the JSON keys Flask returns.

From curl to Code: Where Did Each Line Go?

```
curl --request POST ¥
[1] --url
https://api.ai.sakura.ad.jp/v1/documents/chat/ ¥
    --header 'Accept: application/json' ¥
[2] --header "Authorization: Bearer
${AI_ENGINE_TOKEN}" ¥
    --header 'Content-Type: application/json' ¥
[9] --data '{
[3]   "model": "multilingual-e5-large",
[4]   "chat_model": "gpt-oss-120b",
[5]   "query": "Explain KOSEN-KMITL and Monozukuri.",
[6]   "top_k": 3,
[7]   "threshold": 0.6,
[8]   "distance_type": "cosine"
}'
```

	Where it maps in app.py
1	AI_ENGINE_URL (constant at the top of the file)
2	token = os.getenv(TOKEN_ENV_NAME) → headers["Authorization"]
3	EMBEDDING_MODEL(constant, fixed value)
4	CHAT_MODEL(constant, fixed value)
5	The query variable sent from the web form (blank __ (1) __)
6	The top_k variable sent from the web form (blank __ (2) __)
7	The threshold variable sent from the web form (blank __ (3) __)
8	DISTANCE_TYPE(constant, fixed value)
9	The payload dict returned by build_payload() → requests.post(..., json=__ (5) __)

Wrapping Up the Blanks

Once all 7 are filled in, save and move on.

✓ app.py build_payload: **query** / **top_k** / **threshold**

✓ app.py requests.post: **AI_ENGINE_URL** / **payload**

● Don't forget to save. A missed blank causes an error on the next run.

Once you're done, on to 6 Run It.

Hands-on 3

Run It

Start the server, ask it a question from the browser, and finally shut the server down.

Running Flow

From start to finish, we'll go through these 4 steps.

- 1 `(venv)` enter it (run `activate` if it's not showing)
- 2 `python app.py` to start the server
- 3 Open the displayed URL in your browser
- 4 Shut it down with `Ctrl + C`

Start it up, use it, then shut it down. That's one full cycle.

Start the Server

(**venv**): check the token, then start it.

```
(venv) $ echo $AI_ENGINE_TOKEN    # check the token  
<OK if a token appears here>
```

```
(venv) $ python app.py            # start the server
```

- If (**venv**) isn't showing, enter it with `source venv/bin/activate`.
- If the token is empty, run `export AI_ENGINE_TOKEN=...` first.

Once you're ready, `python app.py` starts the server.

What You'll See on Startup

Once it starts, the terminal shows something like this.

```
* Serving Flask app 'app'
* Debug mode: on
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
```

- Running on ... — that line is your app's URL.
- It's debug mode, so saving your code reloads it automatically.
- Press CTRL+C to quit — that's how to stop it.

Open this URL in your browser, and the app is ready to use.

Why 127.0.0.1 : 5000?

Let's break down what the displayed URL means.

`http://127.0.0.1:5000`

127.0.0.1

Stays inside your own PC. Never exposed externally or to the internet (safe for classroom use).

5000

Flask's default dev-server port. The number that serves as the app's entry point.

- Specified in `app.py`'s `app.run(host="127.0.0.1", port=5000)`.

A safe development server that only runs inside your own PC.

Open It in Your Browser

Open the displayed URL and actually ask it a question.

- In the terminal, **<http://127.0.0.1:5000>** — **Ctrl + click**
(if it doesn't become a link, type it into the browser's address bar by hand)
- Once the page opens, type your question and set top_k / threshold.
- Press “Ask” and the answer and its sources appear.
 - ✓ Once the page opens, ask something about the documents you registered in Module 1 — and set top_k / threshold.
 - ✓ Tip: match your question's language to your registered documents for the best results.

Along with the answer, you can also check the sources (name, chunk number, distance, content).

The screenshot shows a web application titled "RAG Q&A App". Below the title, there is a description: "Sakura AI Engine answers your question by referencing the registered documents. This app sends what you type on the Web screen to Flask, and Flask calls the Sakura AI Engine API." The interface has two main sections. The first section, labeled "Question", contains a text input field with the placeholder text "Enter your question about the registered documents" and a small icon in the bottom right corner. The second section, labeled "RAG Parameters", contains two controls. The first is "top_k", which has a text input field with the value "3" and a description: "Specifies how many of the top, most relevant chunks to retrieve for the question. A larger value references more chunks." The second is "threshold", which has a horizontal slider bar and a description: "Specifies the threshold used when searching documents." The slider bar is currently set to 0.3, and the text "Current value: 0.3" is displayed to the right of the slider.

Nice work! You just called AI Engine through your own Web app, start to finish.

Try It Out: What Happens If...?

Now that it's working, let's experiment a bit!

- **Ask about something not in your registered documents**

What happens: the answer likely says it can't find relevant information — a good way to confirm the app is really searching your documents, not just guessing.

- **Register a new document from the Sakura AI Engine Control Panel, then ask about it**

What happens: check whether the new document gets pulled in as a source — a live look at how documents_chat searches everything you've registered. registered.

- **Ask in a different language than your registered documents**

What happens: search quality may vary — multilingual-e5-large supports cross-lingual search, but matching your question's language to your documents usually works best.

There's no single right answer here. The point is to see how top_k, threshold, and document language actually change what comes back.

RAG Q&A App

Sakura AI Engine answers your question by referencing the registered documents.

This app sends what you type on the Web screen to Flask, and Flask calls the Sakura AI Engine API.

Question

Enter your question about the registered documents

RAG Parameters

top_k

3

Specifies how many of the top, most relevant chunks to retrieve for the question. A larger value references more chunks.

threshold

Current value: 0.3

Specifies the threshold used when searching documents.

Shut Down the Server

Once you're done, stop the server from the terminal.

```
(venv) $ python app.py
* Running on http://127.0.0.1:5000
^C          # press Ctrl + C
(venv) $
```

- Press **Ctrl + C** in the terminal to stop it. The prompt comes back.
- To run it again, use `python app.py`.
- To leave the virtual env, use `deactivate` (**(venv)** disappears).

Start with `python app.py`, stop with **Ctrl + C. That's the basic operation.**

Optional | For Those Who Finish Early: Extra Challenge 1

Swap the Answer-Generation Model

Switch to a different answer-generation model and compare how the answers differ.

```
CHAT_MODEL = "gpt-oss-120b"           # until now
      ↓
CHAT_MODEL = (try different models and see what works well)  # after (example)
```

Recommended Models

- [preview/gemma-4-31B-it](#): Google's multilingual model.
Its English is natural too, and the difference in style from gpt-oss-120b (general-purpose, larger) is easy to spot.
- * All of these are public previews (may change or end without notice).
- * Copy the model name exactly as shown in the model list / Playground, including the preview/ prefix.

⚠ Never Change This

EMBEDDING_MODEL(multilingual-e5-large) **stays the same.**

Registered documents are already embedded with this model. Change the embedding model and search stops matching up. Only the answer-generation side (**CHAT_MODEL**) is safe to change.

IMPORTANT

Only change CHAT_MODEL. Never change the embedding model!

Optional | For Those Who Finish Early: Extra Challenge 2

Change the Defaults and Logging

Make small changes to the screen and server and see how the behavior changes.

Change the Defaults / Range

- In `index.html`, change the default `top_k` / `threshold` values.
- To widen the upper limit, change `index.html`'s max slider together with `app.py`'s `TOP_K_MAX` (e.g. 10→20).
- Change the number of references or the threshold, and watch how the answer and sources shift.
- * Update both the screen (`index.html`) and the server (`TOP_K_MAX`) together, or validation will reject it.

Add More Backend Logging

`app.py`: add more `print` calls to send data to the terminal.

```
print(f"sources={len(result['sources'])} items")
print(f"answer={result['answer'][:40]}")
return jsonify(result), 200
```

See what's happening on the server side, not just on screen.

Tweak both the screen and the server a bit, and see how the behavior differs.

Add a Feature to Read the Answer Aloud

It takes real work, but it's doable with what you learned today plus a bit of extra knowledge.

Browser → (answer text) → **Flask** → (text) → **AI Engine text-to-speech** → (audio WAV) → **Flask** → (play) → **Browser**

- 1 `index.html`: add a “Read Aloud” button.
- 2 `app.js`: on click, send the currently displayed answer text to Flask's new endpoint `/api/speak`.
- 3 `app.py`: add the route `/api/speak` and call AI Engine's text-to-speech API with requests (token via env var again).
- 4 `app.py`: return the audio (WAV) that comes back to the browser.
- 5 `app.js`: play the received audio → save and re-run to check the read-aloud works.

Same as with RAG, “browser → Flask → AI Engine” — try building it this time for speech synthesis.

If you're not confident with programming, you could even have AI Engine help you code the implementation...

Hints for the Read-Aloud Feature

OpenAI TTS-compatible. Send `model` ▪ `voice` ▪ `input` and get audio back.

`app.py` (new route / also import `Response`)

```
SPEECH_URL = "https://api.ai.sakura.ad.jp
              /v1/audio/speech"

@app.route("/api/speak", methods=["POST"])
def speak():
    text = request.get_json().get("text", "")
    payload = {"model": "<speech model ID>",
              "voice": "<voice ID>", "input": text}
    headers = {"Authorization":
              f"Bearer {os.getenv(TOKEN_ENV_NAME)}"}
    r = requests.post(SPEECH_URL,
                      headers=headers, json=payload, timeout=60)
    return Response(r.content,
                    mimetype="audio/wav")
```

`static/app.js` (send the answer and play the audio)

```
speakBtn.addEventListener("click",
  async () => {
    const res = await fetch("/api/speak", {
      method: "POST",
      headers: {"Content-Type":
                "application/json"},
      body: JSON.stringify({
        text: answerText.textContent })
    });
    const blob = await res.blob();
    const url = URL.createObjectURL(blob);
    new Audio(url).play();
  });
```

- Check the speech model ID and voice ID under “Available Speech Models” in the Control Panel (terms agreement required first).
Example: `model`=“kasukabetsumugi” / `voice`=“normal” (here, Kasukabe Tsumugi’s voice). Other voices, like Zundamon, are available too.
- For the latest endpoint and parameter details, check AI Engine’s API documentation.

AI Engine’s text-to-speech API. Think about whether you can build it with the same flow as RAG.

Today's Wrap-Up

Great work!! Let's check off what you can now do.



curl used to call AI Engine — now you can call it from a Web app



You experienced the flow: Browser → JavaScript → Flask → AI Engine.



curl and requests — you learned to read the mapping between them



You filled in the blanks in the code, ran it, and confirmed top_k / threshold and the sources



You saw why the token stays hidden on the Flask side, and how it's implemented